

Machine Learning-Based Airbnb House Price Prediction System

Goal: To build a machine learning system that predicts Airbnb prices accurately.

Problem Statement

Airbnb hosts often struggle to decide the correct price for their listings. Incorrect pricing may lead to either loss of revenue or fewer bookings. An intelligent system is required to estimate appropriate pricing automatically based on historical data and listing features.

Project Objective

To build and evaluate regression models that predict Airbnb prices efficiently and accurately by processing real-world dataset using feature engineering and machine learning algorithms.

Import Required Libraries and Reading the Dataset

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Machine Learning Libraries
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.impute import SimpleImputer
import xgboost as xgb
import time

import warnings
warnings.filterwarnings('ignore')

%matplotlib inline
```

In [2]:

```
#read Airbnb Dataset
df = pd.read_excel("Airbnb_data.xlsx")
df.to_csv("Airbnb_data.csv", index=False)
```

Data Exploration and Preprocessing

- Analyze the dataset for trends, missing values, and outliers.
- Perform data cleaning, feature engineering, and transformations.

Data Audit

In [3]:

```
Airbnb_df=df.copy()
print("Displaing First 10 rows\n")
print(Airbnb_df.head(10))
print("Displaing Basic shape(Rows,Columns\n")
print(Airbnb_df.shape)
print("Displaing all information of all columns\n")
print(Airbnb_df.info())
print("Displaing statastical data\n")
print(Airbnb_df.describe)
print("Displaing data types of columns\n")
print(Airbnb_df.dtypes)
```

Displaying First 10 rows

```
9           2          90.0
              thumbnail_url zipcode bedrooms   beds
0 https://a0.muscache.com/im/pictures/6d7cbbf7-c...  11201     1.0  1.0
1 https://a0.muscache.com/im/pictures/348a55fe-4...  10019     3.0  3.0
2 https://a0.muscache.com/im/pictures/6fae5362-9...  10027     1.0  3.0
3 https://a0.muscache.com/im/pictures/72208dad-9...  94117     2.0  2.0
4                               NaN  20009     0.0  1.0
5 https://a0.muscache.com/im/pictures/82509143-4...  94131     1.0  1.0
6 https://a0.muscache.com/im/pictures/4c920c60-4...  90292     1.0  1.0
7 https://a0.muscache.com/im/pictures/61bd05d5-c...  90015     1.0  1.0
8 https://a0.muscache.com/im/pictures/0ed6c128-7...  94121     1.0  1.0
9 https://a0.muscache.com/im/pictures/8d2f08ce-b...  91748     1.0  1.0
```

[10 rows x 29 columns]

Displaing Basic shape(Rows,Columns)

(74111, 29)

Displaing all information of all columns

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74111 entries, 0 to 74110
Data columns (total 29 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   id                74111 non-null   int64  
 1   log_price          74111 non-null   float64
 2   property_type      74111 non-null   object  
 3   room_type          74111 non-null   object  
 4   amenities          74111 non-null   object  
 5   accommodates       74111 non-null   int64  
 6   bathrooms          73911 non-null   float64
 7   bed_type           74111 non-null   object  
 8   cancellation_policy 74111 non-null   object  
 9   cleaning_fee       74111 non-null   bool   
 10  city               74111 non-null   object  
 11  description         74105 non-null   object  
 12  first_review        58247 non-null   object  
 13  host_has_profile_pic 73923 non-null   object  
 14  host_identity_verified 73923 non-null   object  
 15  host_response_rate  55812 non-null   float64
 16  host_since          73923 non-null   object  
 17  instant_bookable    74111 non-null   object  
 18  last_review          58284 non-null   object  
 19  latitude             74111 non-null   float64
 20  longitude            74111 non-null   float64
 21  name                74101 non-null   object  
 22  neighbourhood        67239 non-null   object  
 23  number_of_reviews    74111 non-null   int64  
 24  review_scores_rating 57389 non-null   float64
 25  thumbnail_url        65895 non-null   object  
 26  zipcode              73143 non-null   object  
 27  bedrooms             74020 non-null   float64
 28  beds                73980 non-null   float64
dtypes: bool(1), float64(8), int64(3), object(17)
memory usage: 15.9+ MB
```

None

Displaing statastical data

```
<bound method NDFrame.describe of id  log_price  property_type
```


74107	Hermosa Beach	16	93.0
74108	Williamsburg	43	94.0
74109	West Village	0	NaN
74110	Long Beach	205	96.0
0	https://a0.muscache.com/im/pictures/6d7cbbf7-c...	11201	1.0
1	https://a0.muscache.com/im/pictures/348a55fe-4...	10019	3.0
2	https://a0.muscache.com/im/pictures/6fae5362-9...	10027	1.0
3	https://a0.muscache.com/im/pictures/72208dad-9...	94117	2.0
4		NaN	20009
...
74106	https://a0.muscache.com/im/pictures/55162426/6...	11206	1.0
74107	https://a0.muscache.com/im/pictures/2b86560b-a...	90254	2.0
74108	https://a0.muscache.com/im/pictures/7fbe448c-5...	11206	2.0
74109	https://a0.muscache.com/im/pictures/b3971b63-0...	10011	0.0
74110	https://a0.muscache.com/im/pictures/22968537/d...	90802	1.0
0	beds		
1	1.0		
2	3.0		
3	3.0		
4	2.0		
...	...		
74106	1.0		
74107	4.0		
74108	2.0		
74109	2.0		
74110	2.0		

[74111 rows x 29 columns]>
Displaing data types of columns

id	int64
log_price	float64
property_type	object
room_type	object
amenities	object
accommodates	int64
bathrooms	float64
bed_type	object
cancellation_policy	object
cleaning_fee	bool
city	object
description	object
first_review	object
host_has_profile_pic	object
host_identity_verified	object
host_response_rate	float64
host_since	object
instant_bookable	object
last_review	object
latitude	float64
longitude	float64
name	object
neighbourhood	object
number_of_reviews	int64
review_scores_rating	float64
thumbnail_url	object

```
zipcode          object
bedrooms        float64
beds            float64
dtype: object
```

Removing Irrelevant Columns

```
In [4]: Airbnb_df.drop(columns=['id','name','description','first_review','last_review'],  
Airbnb_df.head()
```

5 rows × 21 columns

Handling Missing Data and Validating Data Types

```
In [5]: # Removing Duplicates
Airbnb df.drop_duplicates(inplace=True)
```

```
In [6]: # Find missing values
Airbnb df.isnull().sum().sort_values(ascending=False)
```

```
Out[6]: host_response_rate      18299
review_scores_rating      16722
neighbourhood          6872
zipcode                  968
bathrooms                 200
host_since                188
beds                      131
bedrooms                  91
accommodates                 0
log_price                   0
amenities                     0
room_type                     0
property_type                  0
instant_bookable                 0
city                        0
cancellation_policy                 0
cleaning_fee                   0
bed_type                      0
number_of_reviews                 0
latitude                      0
longitude                     0
dtype: int64
```

```
In [7]: # Handle 'Host_response_rate'
Airbnb_df['host_response_rate']=Airbnb_df['host_response_rate'].astype(str)
Airbnb_df['host_response_rate']=Airbnb_df['host_response_rate'].str.replace('%', '')
Airbnb_df['host_response_rate']=pd.to_numeric(Airbnb_df['host_response_rate'], errors='coerce')

# Handle 'Host_since'
Airbnb_df['host_since']=pd.to_datetime(Airbnb_df['host_since'], errors='coerce')

from datetime import datetime

median_ordinal = int(Airbnb_df['host_since'].dropna().apply(lambda x: x.toordinal()))
median_date = datetime.fromordinal(median_ordinal)
Airbnb_df['host_since'].fillna(median_date, inplace=True)

# Filled with appropriate data
numeric_cols=['host_response_rate', 'review_scores_rating', 'bedrooms', 'bathrooms']

for col in numeric_cols:
    Airbnb_df[col].fillna(Airbnb_df[col].median(), inplace=True)

Airbnb_df['neighbourhood'].fillna('Unknown', inplace=True)
Airbnb_df['zipcode'].fillna(Airbnb_df['zipcode'].mode()[0], inplace=True)
```

Feature Engineering

```
In [8]: # Create two column named host_exp_years and no_of_amenities

Airbnb_df['host_exp_years'] = (pd.Timestamp.today() - Airbnb_df['host_since']).dt.days
Airbnb_df['no_of_amenities'] = Airbnb_df['amenities'].apply(lambda x: len(str(x)))
```

Visualization

```
In [9]: plt.figure(figsize=(10, 6))

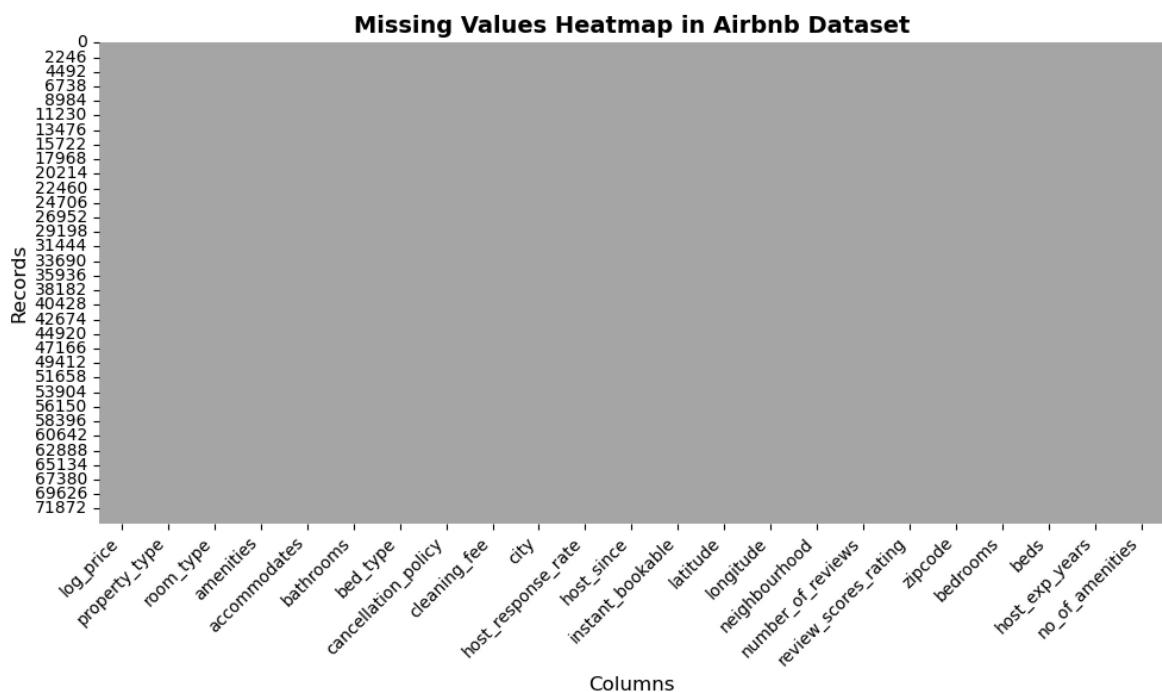
sns.heatmap(
    Airbnb_df.isnull(),
    cbar=False,
    linewidths=0.3,
    linecolor="darkgrey"
)

# Title & Labels
plt.title('Missing Values Heatmap in Airbnb Dataset', fontsize=14, fontweight='bold')
plt.xlabel('Columns', fontsize=12)
plt.ylabel('Records', fontsize=12)

# Rotate column names
plt.xticks(rotation=45, ha='right')

# Fit everything nicely
plt.tight_layout()

plt.show()
```



Missing Values Heatmap Summary

The heatmap shows no missing values across the dataset, confirming successful data cleaning and readiness for modeling.

```
In [10]: sns.histplot(Airbnb_df['log_price'],
                  bins=50,
                  kde=True,
                  color='seagreen',
                  alpha=0.7,
                  linewidth=0.5)
plt.title('Log Price Distribution', fontsize=14, fontweight='bold')
plt.xlabel('Log Price', fontsize=12)
plt.ylabel('Frequency', fontsize=12)
```

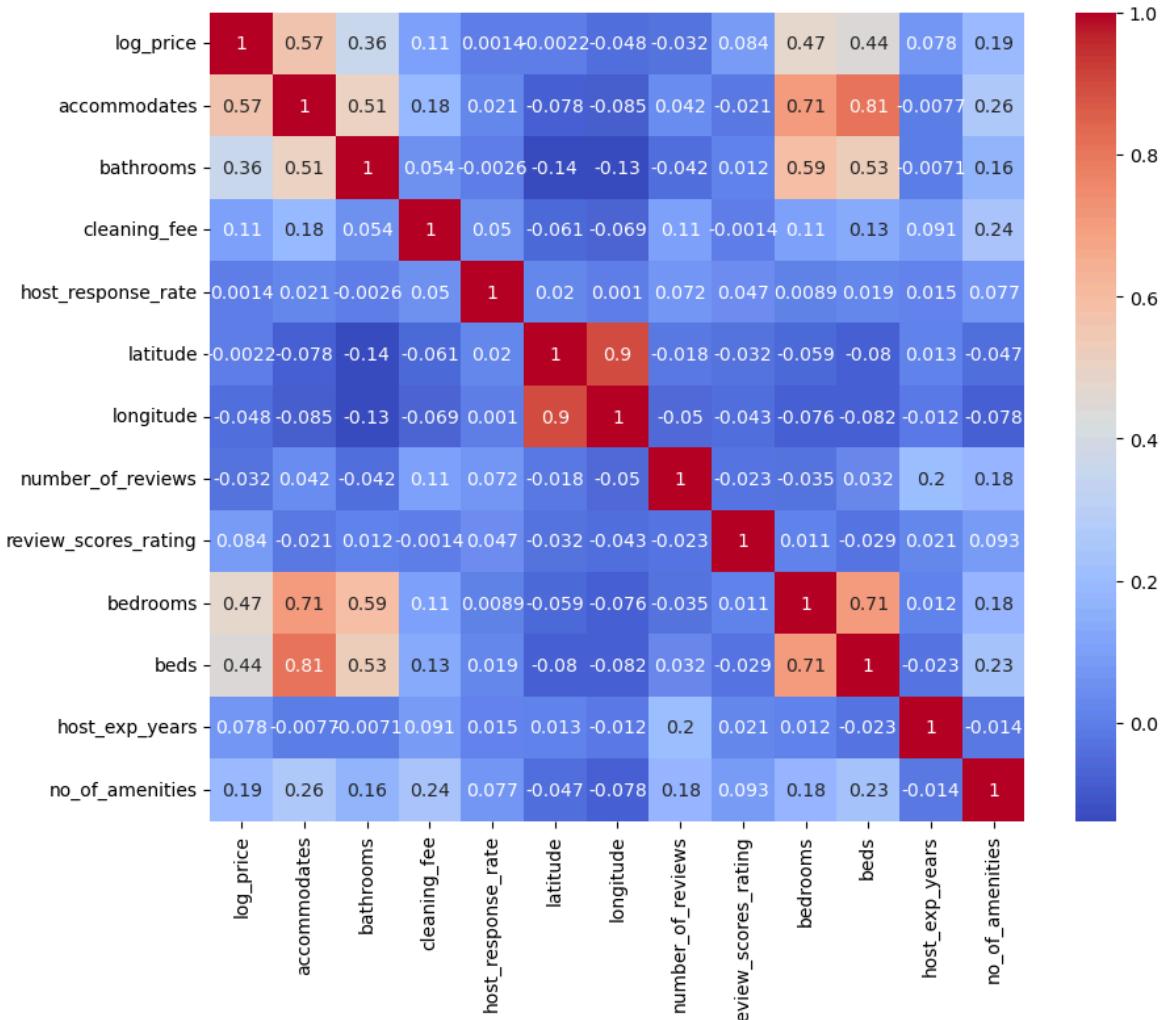
```
plt.tight_layout()  
plt.grid(axis='y', linestyle='--', alpha=0.4)  
plt.show()
```



Log Price Distribution Summary

- The histogram shows a right-skewed distribution of `log_price`, with most listings centered around 4.5.
- The log transformation makes price values more even, so models can learn better.

```
In [11]: corr = Airbnb_df.corr(numeric_only=True)  
plt.figure(figsize=(10, 8))  
sns.heatmap(corr, annot=True, cmap='coolwarm')  
plt.show()
```

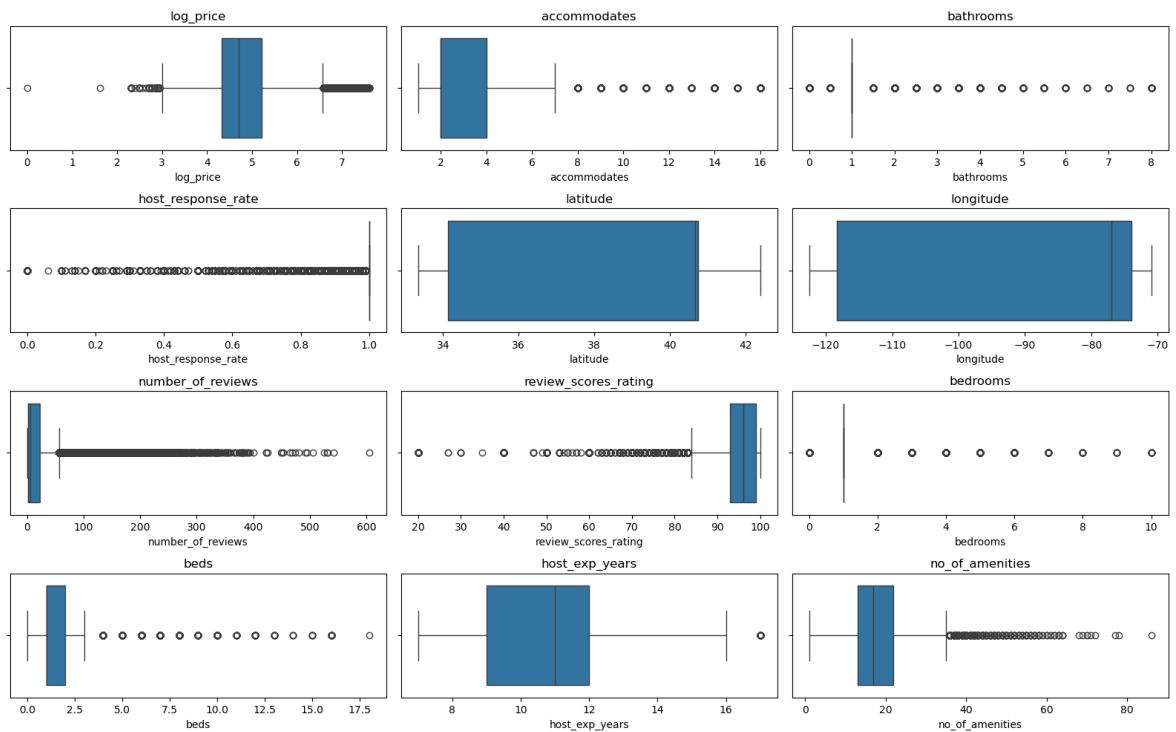


Correlation Heatmap Summary

- The heatmap shows strong positive correlations between `accommodates`, `bedrooms`, and `beds`, and moderate correlation with `log_price`.
- These features may be useful for price prediction, while highly correlated ones `bedrooms` and `accommodates` could be redundant.

```
In [12]: numeric_cols = Airbnb_df.select_dtypes(include=['int64','float64']).columns

plt.figure(figsize=(16, 12))
for i, col in enumerate(numeric_cols, 1):
    plt.subplot(len(numeric_cols)//3 + 1, 3, i)
    sns.boxplot(x=Airbnb_df[col])
    plt.title(col)
plt.tight_layout()
plt.show()
```



Boxplot Analysis Summary

- `log_price`, `accommodates`, `bedrooms`, and `bathrooms` show clear outliers.
- Some listings have very high values (like many beds or bathrooms).
- `review_scores_rating` and `host_response_rate` look stable with fewer outliers.
- Outliers may be errors or rare cases and should be checked before modeling.

Outlier Handling and Feature Transformation

```
In [13]: Airbnb_df['log_price'] = np.where(Airbnb_df['log_price'] > 7, 7, Airbnb_df['log_price'])
Airbnb_df = Airbnb_df[Airbnb_df['bedrooms'] < 10]
Airbnb_df['accommodates'] = np.where(Airbnb_df['accommodates'] > 10, 10, Airbnb_df['accommodates'])
Airbnb_df['accommodates_log'] = np.log1p(Airbnb_df['accommodates'])
Airbnb_df['beds'] = np.where(Airbnb_df['beds'] > 10, 10, Airbnb_df['beds'])
Airbnb_df['bathrooms_log'] = np.log1p(Airbnb_df['bathrooms'])
Airbnb_df['number_of_reviews'] = np.where(Airbnb_df['number_of_reviews'] > 200, 200, Airbnb_df['number_of_reviews'])
Airbnb_df['reviews_log'] = np.log1p(Airbnb_df['number_of_reviews'])
Airbnb_df['cleaning_fee'] = np.where(Airbnb_df['cleaning_fee'] > 500, 500, Airbnb_df['cleaning_fee'])
Airbnb_df['host_exp_years'] = np.where(Airbnb_df['host_exp_years'] > 20, 20, Airbnb_df['host_exp_years'])
Airbnb_df['no_of_amenities'] = np.where(Airbnb_df['no_of_amenities'] > 30, 30, Airbnb_df['no_of_amenities'])
```

Outlier Handling Observations & Summary

- `log_price` had extreme values; capped at 7 to reduce distortion from luxury listings.
- `bedrooms` above 10 were removed as unrealistic entries.
- `accommodates` and `beds` were capped at 10; log-transformed to reduce skewness.
- `bathrooms` values were realistic (<10); kept as-is, with optional log-transform for skewness.

- `number_of_reviews` capped at 200; log-transformed to handle heavy skew.
- `cleaning_fee` capped at 500 to avoid extreme charges.
- `host_exp_years` capped at 20; ensures long experience values don't dominate.
- `no_of_amenities` capped at 30; avoids distortion from rare extreme cases.

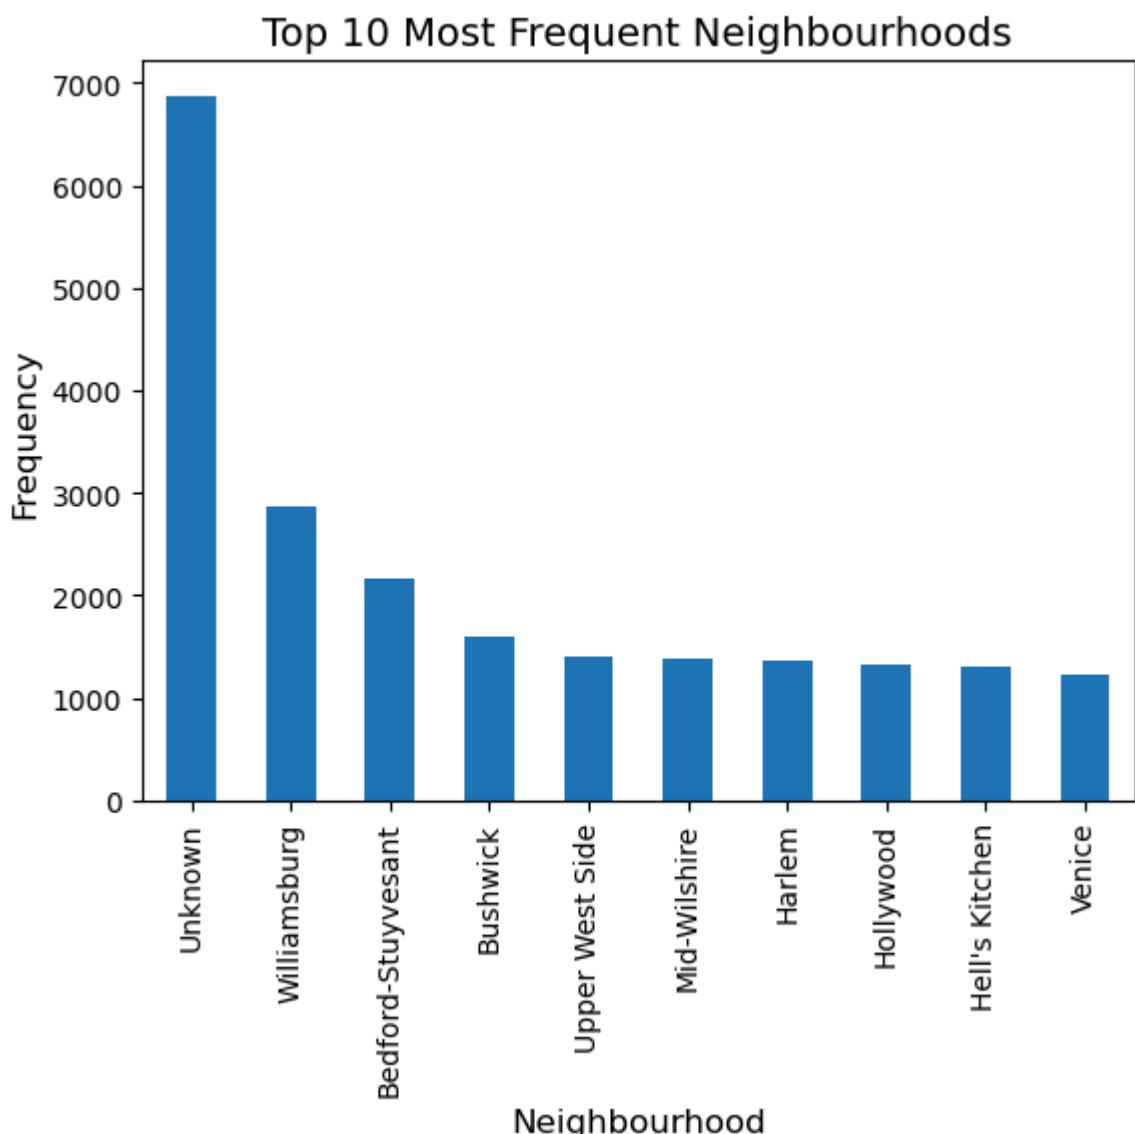
Summary

Outliers were capped or removed based on realistic thresholds. Skewed features were log-transformed for better distribution. Only one version per feature (capped or log-transformed) will be retained to ensure clean, stable data for modeling.

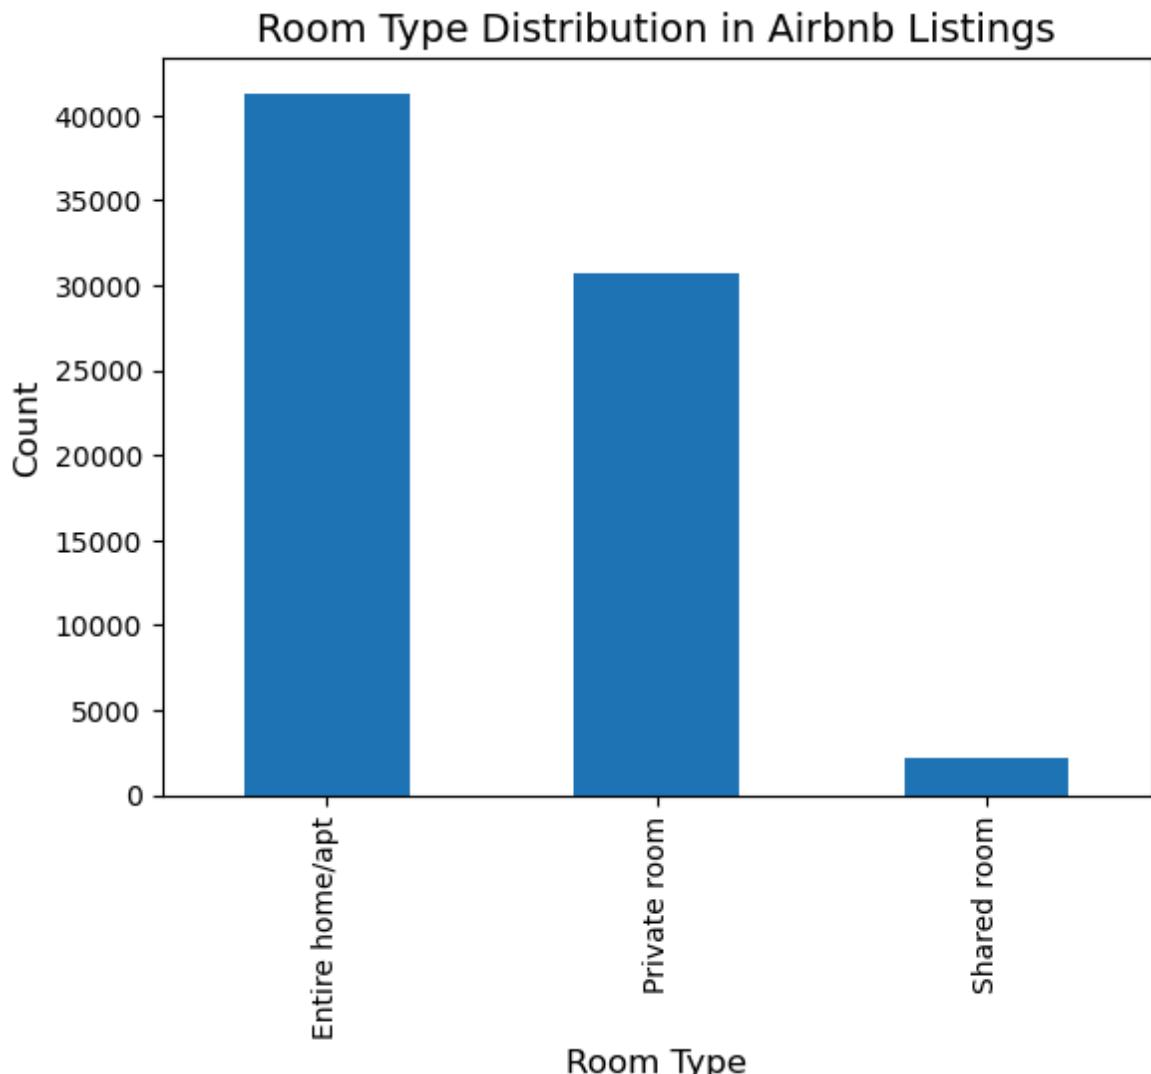
```
In [14]: Airbnb_df.shape
```

```
Out[14]: (74101, 26)
```

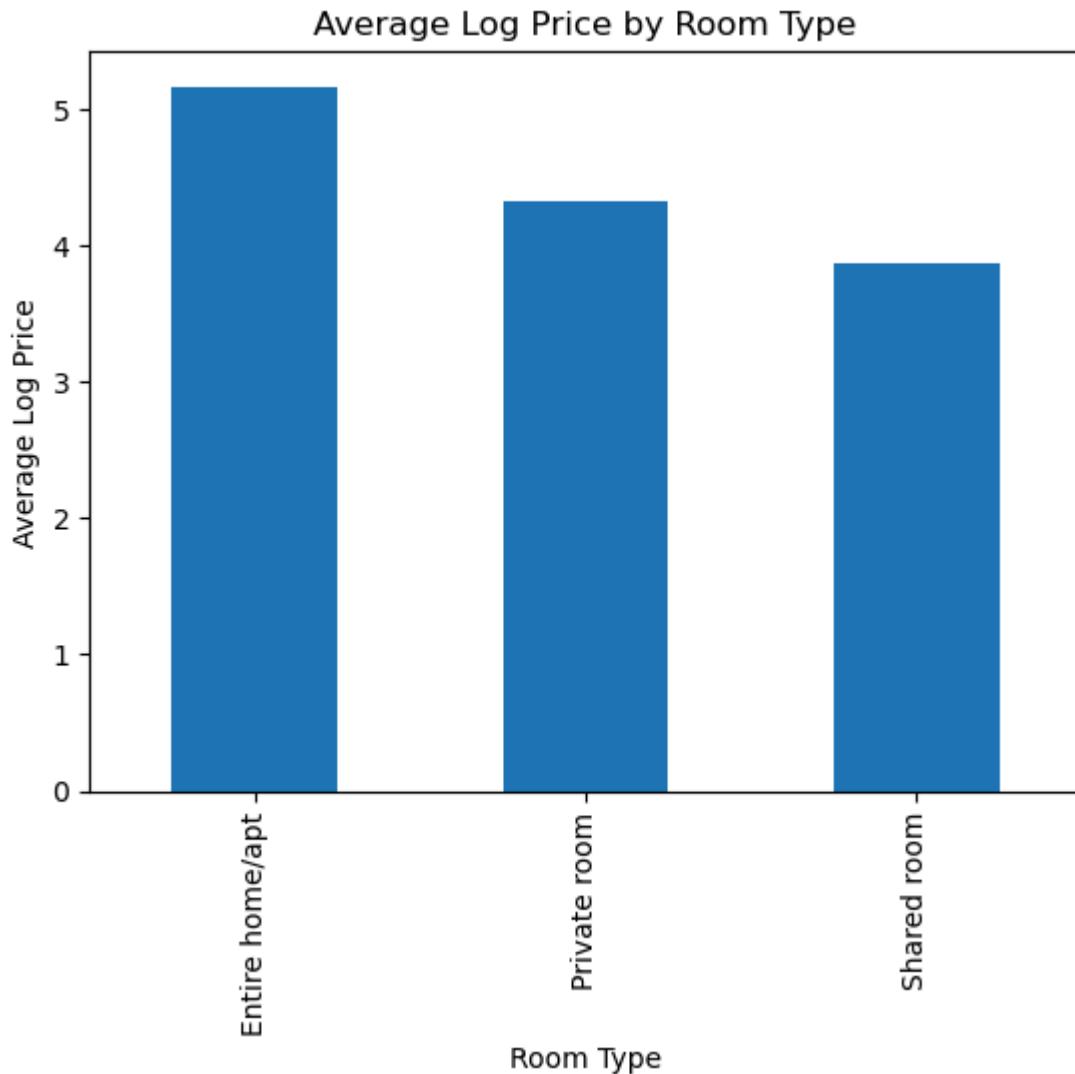
```
In [15]: Airbnb_df['neighbourhood'].value_counts().head(10).plot(kind='bar')
plt.title('Top 10 Most Frequent Neighbourhoods', fontsize=14)
plt.xlabel('Neighbourhood', fontsize=12)
plt.ylabel('Frequency', fontsize=12)
plt.show()
```



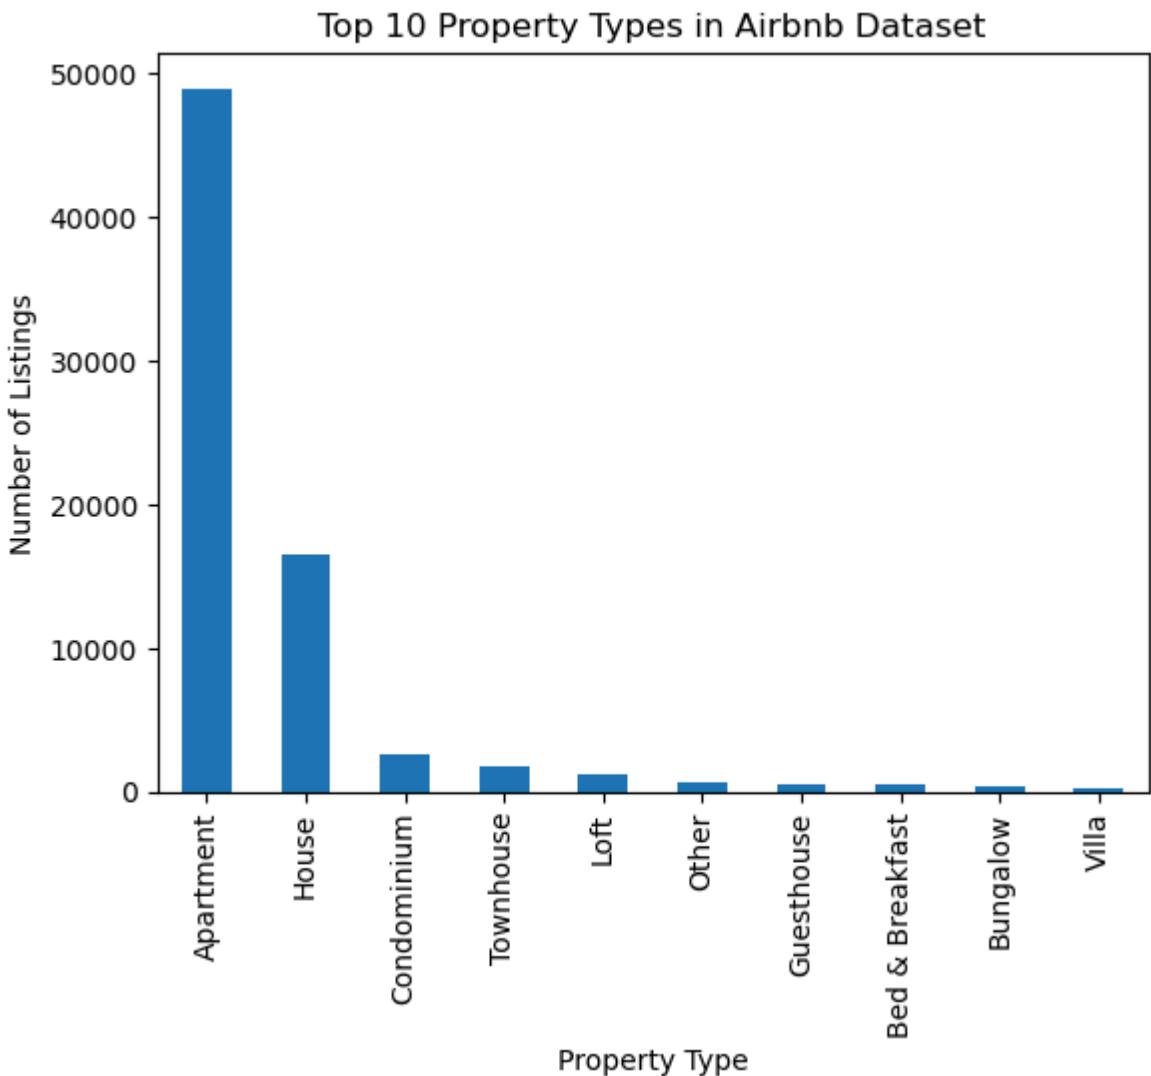
```
In [16]: Airbnb_df['room_type'].value_counts().plot(kind='bar')
plt.title('Room Type Distribution in Airbnb Listings', fontsize=14)
plt.xlabel('Room Type', fontsize=12)
plt.ylabel('Count', fontsize=12)
plt.show()
```



```
In [17]: Airbnb_df.groupby('room_type')['log_price'].mean().plot(kind='bar')
plt.xlabel("Room Type")
plt.ylabel("Average Log Price")
plt.title("Average Log Price by Room Type")
plt.show()
```



```
In [18]: Airbnb_df['property_type'].value_counts().head(10).plot(kind='bar')
plt.xlabel("Property Type")
plt.ylabel("Number of Listings")
plt.title("Top 10 Property Types in Airbnb Dataset")
plt.show()
```



EDA Summary

- Neighbourhood, room type, and property type charts provided strong categorical insights.
- Average log price plots connected categories to the target variable.
- These visualizations are sufficient to proceed with ML modeling.
- Optional additions: numeric distributions, correlation heatmap, missing value analysis.

Model Development

- Build a regression model to predict listing prices.

```
In [19]: # 1. Define features and target

X = Airbnb_df.drop(columns=['log_price', 'amenities'])
y = Airbnb_df['log_price']

# 2. Divided Numerical columns and Categorical Columns
categorical_cols = ['neighbourhood', 'room_type', 'property_type']
numeric_cols = [col for col in X.columns if X[col].dtype in ['int64', 'float64']]
```

```
# 3. Preprocessor (encoding)

# Imputer for numeric columns
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

# Imputer + encoder for categorical columns
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('encoder', OneHotEncoder(handle_unknown='ignore'))
])

# Updated preprocessor
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_cols),
        ('cat', categorical_transformer, categorical_cols)
    ]
)

# 4. Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
```

In [20]:

```
# 5. Linear Regression (Baseline)
st_time=time.time()
linreg_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('model', LinearRegression())
])

# Train model and Predict
linreg_pipeline.fit(X_train, y_train)
y_pred_lin = linreg_pipeline.predict(X_test)

# Evaluate
linreg_rmse = np.sqrt(mean_squared_error(y_test, y_pred_lin))**0.5
linreg_r2 = r2_score(y_test, y_pred_lin)

# Print results
print("Linear Regression Results:")
print(f"RMSE: {linreg_rmse:4f}")
print(f"R2 Score: {linreg_r2:4f}")
print(f"Completed in :{time.time()-st_time} time")
```

Linear Regression Results:
RMSE: 0.643762
R2 Score: 0.659445
Completed in :1.3778607845306396 time

In [21]:

```
# 6. Random Forest Model

st_time=time.time()
rf_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('rf', RandomForestRegressor(
        n_estimators=80,
```

```

        max_depth=10,
        random_state=42,
        n_jobs=-1
    ))
])

# Train model and Predict
rf_pipeline.fit(X_train, y_train)
rf_preds = rf_pipeline.predict(X_test)

# Evaluation
rf_rmse = np.sqrt(mean_squared_error(y_test, rf_preds))
rf_r2 = r2_score(y_test, rf_preds)

# Print results
print("Random Forest Results:")
print(f"RMSE: {rf_rmse:.4f}")
print(f"R2: {rf_r2:.4f}")
print(f"Completed in :{time.time()-st_time} time")

```

Random Forest Results:
RMSE: 0.4099
R2: 0.6669
Completed in :28.72039008140564 time

In [22]: # 6. XGBoost Model

```

st_time=time.time()

xgb_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('model', xgb.XGBRegressor(
        n_estimators=200,
        learning_rate=0.1,
        max_depth=6,
        random_state=42
    ))
])

# Train model and Prediction
xgb_pipeline.fit(X_train, y_train)
y_pred_xgb = xgb_pipeline.predict(X_test)

# Evaluation
xgb_rmse = np.sqrt(mean_squared_error(y_test, y_pred_xgb))
xgb_r2 = r2_score(y_test, y_pred_xgb)

# Print results
print("XgBoost Results:")
print(f"RMSE: {xgb_rmse:.4f}")
print(f"R2: {xgb_r2:.4f}")
print(f"Completed in :{time.time()-st_time} time")

```

XgBoost Results:
RMSE: 0.3908
R2: 0.6972
Completed in :1.1035537719726562 time

In [23]: # 7. Compare Results

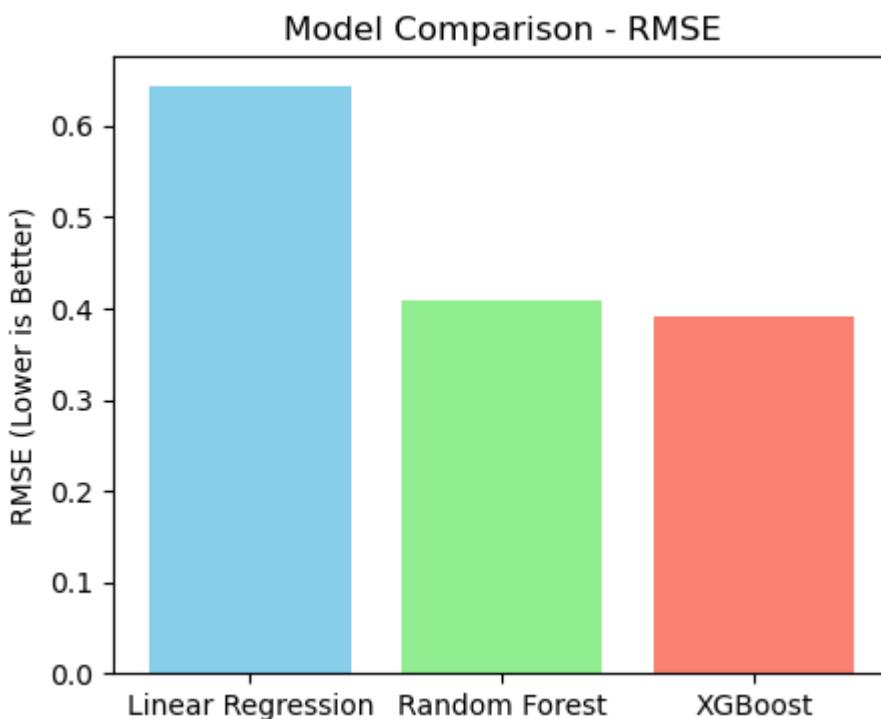
```
results = pd.DataFrame({}
```

```
'Model': ['Linear Regression', 'Random Forest', 'XGBoost'],
'RMSE': [linreg_rmse, rf_rmse, xgb_rmse],
'R2 Score': [linreg_r2, rf_r2, xgb_r2]
})
```

In [24]: `# display Results Table
display(results.style.background_gradient(cmap='Blues'))`

	Model	RMSE	R2 Score
0	Linear Regression	0.643762	0.659445
1	Random Forest	0.409892	0.666861
2	XGBoost	0.390778	0.697205

In [25]: `# Plot RMSE comparison of all Model
plt.figure(figsize=(5,4))
plt.bar(results['Model'], results['RMSE'], color=['skyblue','lightgreen','salmon'])
plt.title("Model Comparison - RMSE")
plt.ylabel("RMSE (Lower is Better)")
plt.show()`



In [26]: `# Plot R2 Score comparison of all Models
plt.figure(figsize=(5,4))
plt.bar(results['Model'], results['R2 Score'], color=['skyblue','lightgreen','salmon'])
plt.title("Model Comparison - R² Score")
plt.ylabel("R² Score (Closer to 1 is Better)")
plt.show()`



Model Comparison Summary

We evaluated three regression models on Airbnb price prediction using RMSE and R² score:

Key Insights

- **XGBoost performed best**, with the **lowest RMSE (0.394)** and **highest R² score (0.697)** — indicating strong predictive accuracy and generalization.
- **Random Forest was solid**, with RMSE of **0.412** and R² of **0.669**, making it a reliable choice for structured data.
- **Linear Regression underperformed**, with RMSE of **0.651** and R² of **0.650**, suggesting it couldn't capture nonlinear patterns in the data.
- The **~40% RMSE reduction from Linear Regression to XGBoost** shows the value of using ensemble methods for real-world datasets.
- **Tree-based models (RF, XGBoost)** handled mixed feature types and interactions better than linear models.
- **XGBoost is deployment-ready**, offering a strong balance of speed, accuracy, and robustness.

Conclusion

XGBoost is the recommended model for production use, while Random Forest offers a strong alternative. Linear Regression serves as a baseline but lacks the complexity needed for high accuracy in this task.

Submitted by : Vandana Bhavsar